

GCS North Star Metric

Account Value Realization Index (AVRI)

"PM is a translation function, between business intent, data reality, and shippable software."

GCS is moving from selling deals to sustaining value.

TCV-based compensation rewarded the signature. The new ARR + consumption hybrid only pays out if customers actually use what they bought. We need ONE metric that balances four very different things.

OLD WORLD: TCV

Book the deal. Bank the credit.

One signal: the contract was signed.

One moment: the day it closed.

One incentive: land it big and walk away.

Blind to: deployment • health • sustained usage

NEW WORLD: ARR + CONSUMPTION

Bookings

Deployment

Tech Health

Sustained Usage

Customer pays only if value is realized. Compensation must reflect the same.

Need: ONE metric that balances all four

The business question I'm answering: ***"What's the health of an account at a point in time, balancing all four dimensions, in a single defensible number?"***

Every existing metric is blind to ≥1 of the four dimensions.

I scored 30+ industry-standard SaaS/CS metrics against the four required dimensions. None covers all four. The closest pattern, Gainsight-style composites, has known gaps for consumption-hybrid businesses.

Metric	Bkgs	Deploy	Tech	Usage	Lead/Lag
TCV (Total Contract Value)	●	○	○	○	Lag
ARR (Annual Recurring Rev)	●	○	○	○	Lag
NRR (Net Revenue Retention)	●	○	○	◐	Lag
Raw Utilization (90d)	○	◐	○	●	Lag
Latest Health Color	○	○	◐	○	Lead
Gainsight CHS (composite)	◐	◐	●	●	Lead
AVRI (proposed)	●	●	●	●	Lead

● covers ◐ partial ○ blind

Build vs Buy: Gainsight CHS is the closest commercial alternative, same composite pattern, but doesn't natively handle commit-vs-consume, trajectory, or the 5 brief-specified edge cases. AVRI is that pattern, specialized for PANW's consumption-hybrid model.

On 766 in-scope accounts:

AVRI (ours)



69% / 8% / 22%

Naive CHS (industry baseline)



59% / 22% / 19%

AVRI surfaces **more risk** than naive CHS, and is **more lenient on the right things** (overage = upsell, not failure).

AVRI: a 0–100 composite that balances all four dimensions.

CR**Commit Realization**

weight 30%

Are they using what they paid for?

90-day consumption ÷ 90-day commit. Piecewise: penalize <50%, plateau 50–110%, gentle decay >150%.

UM**Usage Momentum**

weight 30%

Is consumption holding up?

90-day avg ÷ 12-month baseline. Catches spike-and-drop. Level guard zeros UM if annual usage <5% of commit.

DM**Deployment Maturity**

weight 20%

Is the product genuinely operationalized?

Active days in last 90 ÷ 90. Production version: also license provisioning, feature breadth, milestones.

TH**Technical Health**

weight 20%

Does it work for them?

Exponentially-decayed health color (0.95/day). Production: also Sev-1 freq, SLA attainment, escalations.

$$AVRI = 0.30 \cdot CR + 0.30 \cdot UM + 0.20 \cdot DM + 0.20 \cdot TH$$

Floor rule: if TH < 30, cap AVRI at 50.

GREEN ≥ 75*Healthy · expand***YELLOW 50–74***Watch · CSM action***RED < 50***At risk · escalate*

Three accounts where AVRI gets it right and existing metrics don't.

Inductive validation: the data was generated to be realistic, not to flatter the metric. AVRI's edge cases are visible on real accounts. Then we'll switch to the live dashboard.

S H E L F W A R E	F L O O R R U L E	R E W A R D E D
ACC -00202 Financial Services • \$3.5M ARR Naive CHS:40 (yellow) AVRI:19.5 (red) <i>0 active days in 90. Latest health color was green from earlier in the year, naive CHS is fooled by it. AVRI sees through.</i>	ACC -00226 Technology • \$8M ARR Naive CHS:70 (yellow) AVRI:50 (yellow) <i>Consumption looks great (130% util, all 4 pillars high) BUT TH = 15. Without floor rule AVRI = 86 (green). Rule correctly caps it.</i>	ACC -00876 Manufacturing • \$246K ARR Naive CHS:57 (yellow) AVRI:92 (green) <i>All pillars high; latest health color happened to be red (one bad week). Naive CHS over-penalizes; AVRI's decay-weighted TH gets the trend right.</i>

→ LIVE DEMO Streamlit • At-Risk Renewals tab • \$10.3M across 20 accounts renewing in 90 days, flagged Yellow or Red.

HUMAN IN THE LOOP AVRI scores; dashboard surfaces; CSM acts. The metric directs the human, not replaces them.

SUCCESS CRITERIA (1) % of Red accounts churning in 90d (2) Reduction in Green→Lost surprise renewals (3) Expansion conversion in 110–150% util band

What I learned, what v2 looks like, how to ship it.

Built end-to-end in ~8 hours of focused spec-driven AI work: data generator (Python), BigQuery pipeline (SQL), dashboard (Streamlit), 13 documented pitfalls captured in `lessons_learned.md`. Below: what changes next.

01 What worked + what I'd change

- Specs first, AI codes, the markdown specs are what survive.
- Inductive workflow (data → metrics → refine) beats deductive, stronger narrative for the panel.
- Pre-write validation gates caught structural problems before BigQuery.
- Next time: build CI from day one, would've caught UM zero-baseline bug 2 hours earlier.
- Separate LLMs for code-gen vs code-review.

02 AVRI v2, newer tech

- Agentic data quality: agent watches the pipeline, surfaces anomalies in Slack autonomously.
- LLM-summarized account briefs: per at-risk account, generate CSM-ready talking points from history.
- Tenure-aware cold-start: separate Onboarding RAG state instead of forcing G/Y/R buckets.
- Per-segment weight tuning (Enterprise vs Mid-Market), calibrated against renewal correlation.
- Replace proxies in DM and TH with real signal: license provisioning, Sev-1 frequency, NPS.

03 30/60/90 day rollout

- Day 30, Validate AVRI against historical renewal outcomes; calibrate weights via regression.
- Day 60, Integrate ticket data + license provisioning; replace proxy signals.
- Day 90, Roll into CSM compensation as a 25% factor (alongside book ARR retention).
- Day 180, Dashboard goes from prototype to production (Looker/Tableau for execs; Streamlit retained for diagnostics).
- Day 365, Re-tune weights based on first full year of data.

The repo is reproducible end-to-end. Anyone on your team can clone it and have AVRI scoring real data on day one.